

# Unikernels, Multikernels, Virtual Machine-based Kernels

<http://d3s.mff.cuni.cz-aosy>



CHARLES UNIVERSITY  
Faculty of Mathematics  
and Physics

Department of  
Distributed and  
Dependable  
Systems

D3S

*Martin Děcký*  
decky@d3s.mff.cuni.cz

# Recall: Common OS Taxonomy




- **Special-purpose operating systems**

- Real-time operating systems
- Hypervisors (type 1)
- ...


- **General-purpose operating systems**

- Monolithic kernel
- Single-server microkernel
- Multiserver microkernel
- Hybrid kernel (?)


# Monolithic Kernel



# Single-server Microkernel



# Multiserver Microkernel



# Examples



- **Monolithic kernel**
  - Linux, Solaris (UTS), Windows, FreeBSD, NetBSD, OpenBSD, OpenVMS, MS-DOS, RISC OS
- **Single-server microkernel**
  - CMU Mach, AmigaOS, NeXTSTEP
- **Multiserver microkernel**
  - HelenOS, MINIX 3, Genode, GNU/Hurd, Redox, RefOS (seL4)


# Hybrid Kernel



- **No universally accepted definition**


- “Architecture that somewhat resembles a microkernel, but is not a pure microkernel.”
- Windows NT
  - Internal architecture and communication abstractions inspired by CMU Mach (message passing, ports)
  - Originally user space device drivers
    - Later moved into the kernel, more recently some driver classes in user space again
- macOS (iOS, etc.)
  - The core xnu component actually contains the CMU Mach 3.0 microkernel
  - Native and ported (BSD) system functionality in user space
    - Gradual erosion towards kernel drivers
- BeOS (Haiku), Syllable, DragonFly BSD, Plan 9, NetWare, eComStation

# Windows NT




[1]

# macOS




[2]


# Multiserver Microkernel (reprise)




# Hypervisor (Type 1)



# Common Cloud Deployment



# Unikernel



# Unikernel (2)



- **Library operating system**

- Payload (application) merged with the kernel
  - Kernel component acts as a library providing access to the hardware, threading, file systems, etc.
    - Only necessary functionality
  - Mostly static (single image), but there are dynamic variants
  - Code runs in privileged mode and single address space
    - No mode switches, address space switches
    - Syscalls can be replaced by function calls
    - Isolation/security provided by the underlying hypervisor

## ● Exokernel

- MIT since 1994
- Goal: End-to-end principle
  - Limiting the number of abstractions (compared to monolithic kernels)
  - Limiting the communication complexity (compared to microkernels)
- Co-existence with a regular kernel
- ExOS (MIT)
- Nemesis (University of Cambridge, University of Glasgow, Swedish Institute of Computer Science, Citrix)
  - Multimedia applications

# Unikernel (4)



## ● Rumprun

- POSIX compliant, BSD-compatible run-time environment
  - Original concept: NetBSD anykernel
    - Possibility to compile NetBSD drivers and subsystems either as traditional kernel components or as standalone user space libraries (rump kernels)
    - Rump kernels communicate with the host kernel using syscall interface
  - Replacing the syscall interface of rump kernels with a hypercall interface to the hypervisor
  - “Bare metal” execution also possible

## ● Drawbridge

- Win32-compatible run-time environment
  - Originally a Win32 environment running in a Windows picoprocess

# Unikernel (5)

```
git clone https://github.com/rumpkernel/rumprun.git  
cd rumprun
```

```
git submodule init  
git submodule update
```

(depending on the local C compiler, apply a small patch from  
<https://github.com/rumpkernel/rumprun/issues/86>)

```
export PATH="$PATH:`pwd`/rumprun/bin"
```

```
x86_64-rumprun-netbsd-gcc -o module module.c  
rumprun-bake hw virtio unikernel.bin module
```

```
rumprun kvm -i unikernel.bin
```

# Unikernel (6)



- **OS<sup>v</sup>**


- Linux-compatible application environment
  - Not just a run-time, but a complete OS compatibility
  - Goal: Running hosted applications in unmodified run-time environments for Linux
    - E.g. Java EE application running in Tomcat for Linux
- No notion of users, single address space, processes emulated using threads
  - Similar deployment as in Linux (shell scripting, etc.)




- **Managed language run-time in kernel**

- Clive (Go)
- ClickOS, IncludeOS, HermitCore (C++)
- HaLVM (Haskell)
- LING (Erlang)
- MirageOS (Ocaml)
- Runtime.js (JavaScript)

# Monolithic Kernel (reprise)



# Virtual Machine-based Kernel



# Virtual Machine-based Kernel (2)



## ● Inferno

- Derived from Plan 9 from Bell Labs
  - “Everything is a file” paradigm
    - All global objects represented as file system paths (global namespace)
    - All operations with objects mapped to common file system operations (walk, stat, create, open, read, write, close, etc.)
    - Local object identification using file descriptors
  - Dis virtual machine and Limbo type-safe language
    - Ada-like syntax, Modula inspired modules, concurrency model based on Communicating Sequential Processes, garbage collector
  - File system operations mapped to Styx communication protocol (compatible with 9P2000)
    - Distributed computing


# Virtual Machine-based Kernel (3)



## ● Singularity

- Microsoft (2003 – 2010)
  - Sing# virtual machine
    - Based on Spec#
      - Superset of C# with Eiffel-like specification of code contracts (object invariants, preconditions, postconditions, non-nullable types)
      - Static checker (based on theorem prover)
      - Run-time checker
    - Extends Spec# with support for communication channels and low-level constructs (structures, inline assembler)
    - Bartok just-in-time compiler (CIL to x86)
  - Research prototype
    - Midori expected to be a commercial variant and future replacement of Windows NT (discontinued in 2015)
  - Several similar approaches  
(MOSA, Cosmos for C#; JNode, Phantom OS for Java)

# Multikernel



# Multikernel (2)



## ● Barrelyfish

- ETH Zürich, Microsoft Research
- Side note: Mackerel hardware description language
  - Driver synthesis
- Support for heterogeneous CPU cores
- Common asynchronous messaging abstraction between cores/nodes
  - Explicit inter-core communication (as opposed to cache coherency)
  - Practically no shared memory and state between cores
  - Managing state replicas using distributed algorithms between cores

# Barrelfish: Mackerel Language



- **Describe devices and control registers**

- Used to generate C accessor functions
- Also for other low level data structures

```
device HPET lsbfirst ( addr base ) "High-Precision Event Timer" {
    register gcap_id ro addr(base, 0x0) "General Capabilities and Identification" {
        rev_id             8  "Revision Identification";
        num_tim_cap        5  "Number of Timers";
        count_size_cap     1  "Counter Size";
                           1  mbz;
        leg_rt_cap         1  "Legacy Replacement Rout Capable";
        vendor_id_cap      16 "Vendor ID";
        counter_clk_per_cap 32 "Main Counter Tick Period";
    };
};
```

# Multikernel (3)



## ● Barrelyfish

- Kernel acts as a “CPU driver”
  - Event-driven, single-threaded, non-preemptable
    - Processing syscalls from user space, interrupts from devices and other cores
    - User space processes communicate with the CPU driver using a dispatcher object (local user space thread scheduler)
  - ~10.000 lines of C, ~500 lines of assembler
- Common messaging abstraction does not mean common (sub-optimal) messaging transport
  - Fast path messaging between cache-coherent cores: Sending messages in cache lines
    - Sender writes words sequentially into the cache line (interconnect cache line invalidate)
    - Receiver polls on the last word of the cache line (interconnect cache line fetch)

# Barrelfish: Interface Specification



- **Describe types and messages**

- Used to generate IPC stubs

```
interface timer "Timer service" {  
    // set the one (and only) timeout value (in us) for this client  
    message set_timeout(uint64 timeout);  
    // add the given increment (in us) to the running timer for this client  
    message add_to_timeout(uint64 increment);  
    // cancel the outstanding timeout  
    message cancel_timeout();  
    // wakeup response when the timer is triggered  
    message wakeup();  
    // request for the remaining time of the currently-running timer  
    message get_remaining();  
    // response containing remaining time of running timer  
    message remaining(uint64 time);  
};
```

# Q&A

# References



- [1] Grm wnr, Xyzzy n, [https://commons.wikimedia.org/wiki/File:Windows\\_2000\\_architecture.svg](https://commons.wikimedia.org/wiki/File:Windows_2000_architecture.svg)
- [2] Utente:Sassospicco, [https://commons.wikimedia.org/wiki/File:Diagram\\_of\\_Mac\\_OS\\_X\\_architecture.svg](https://commons.wikimedia.org/wiki/File:Diagram_of_Mac_OS_X_architecture.svg)